


Neural Network




COGNIBOT
AI meets Industry




$$\text{Output of neuron} = Y = f(w1.X1 + w2.X2 + b)$$

Impulses carried toward cell body

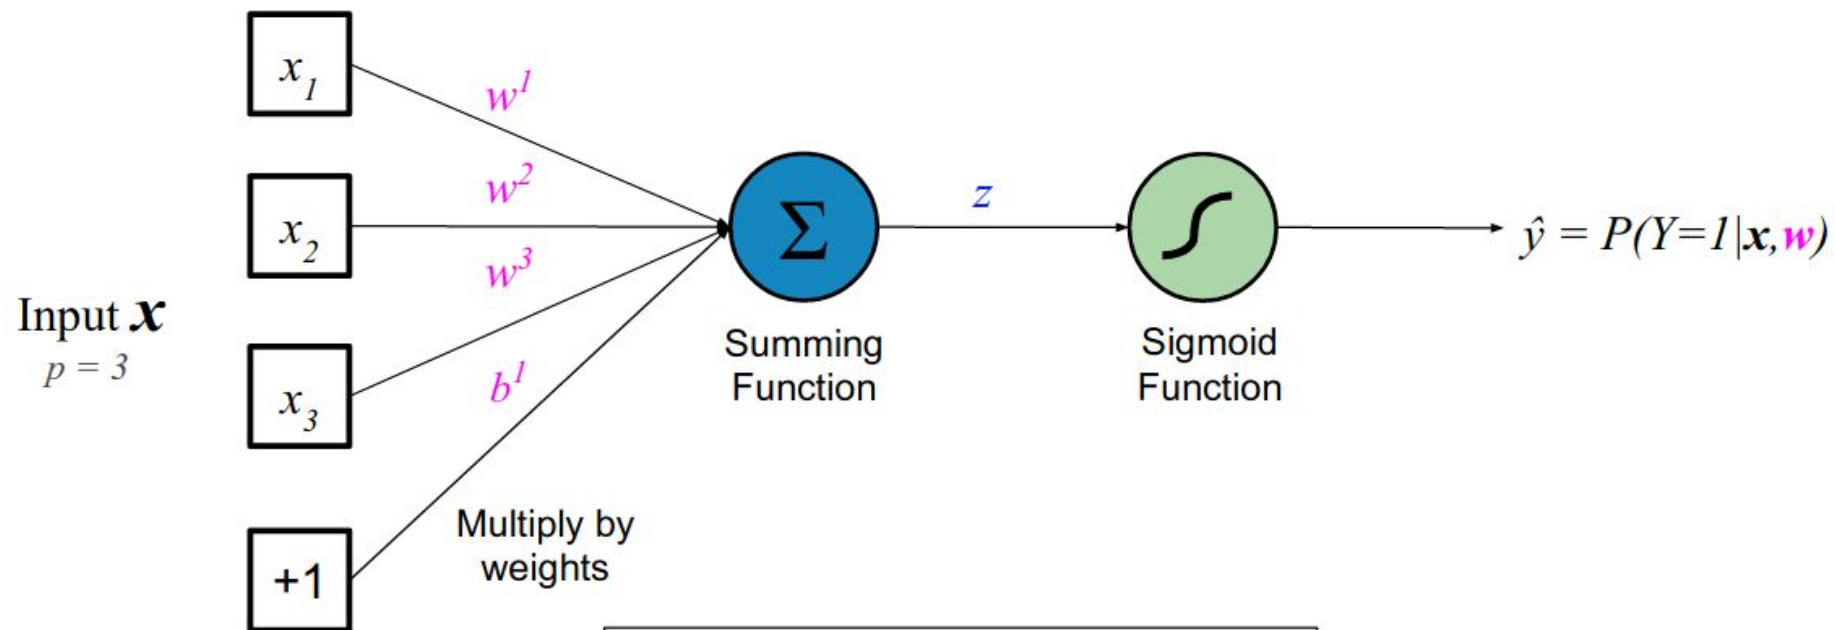


This image by Felipe Perucho
is licensed under [CC-BY 3.0](https://creativecommons.org/licenses/by/3.0/)



NN & Regression


Expanded Logistic Regression



$$\underset{1 \times 1}{z} = \underset{1 \times p}{\mathbf{w}^T} \underset{p \times 1}{\mathbf{x}} + b$$

$$y = \text{sigmoid}(z) = \frac{e^z}{1 + e^z}$$

"Neuron"




1-Layer network



COGNIBOT
AI meets Industry

1-Layer Neural Network (with 4 neurons)



$$\begin{aligned} \mathbf{z} &= \mathbf{W}^T \mathbf{x} \\ \begin{matrix} d \times 1 & d \times p & p \times 1 \end{matrix} \\ \hat{\mathbf{y}} &= \text{sigmoid}(\mathbf{z}) = \frac{e^{\mathbf{z}}}{1 + e^{\mathbf{z}}} \\ \begin{matrix} d \times 1 & d \times 1 \end{matrix} \end{aligned}$$

Element-wise on vector \mathbf{z}

Activation function




COGNIBOT
AI meets Industry

Activation Functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Activation Functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Training



COGNIBOT
AI meets Industry



COGNIBOT
AI meets Industry

See Google Dev.



COGNIBOT
AI meets Industry

Implementation



COGNIBOT
AI meets Industry

Full implementation of training a 2-layer Neural Network needs ~20 lines:

```
1  import numpy as np
2  from numpy.random import randn
3
4  N, D_in, H, D_out = 64, 1000, 100, 10
5  x, y = randn(N, D_in), randn(N, D_out)
6  w1, w2 = randn(D_in, H), randn(H, D_out)
7
8  for t in range(2000):
9      h = 1 / (1 + np.exp(-x.dot(w1)))
10     y_pred = h.dot(w2)
11     loss = np.square(y_pred - y).sum()
12     print(t, loss)
13
14     grad_y_pred = 2.0 * (y_pred - y)
15     grad_w2 = h.T.dot(grad_y_pred)
16     grad_h = grad_y_pred.dot(w2.T)
17     grad_w1 = x.T.dot(grad_h * h * (1 - h))
18
19     w1 -= 1e-4 * grad_w1
20     w2 -= 1e-4 * grad_w2
```

Define the network

Forward pass

Calculate the analytical gradients

Gradient descent